Diabetic Retinopathy Pipeline

Implementation & Session Documentation

Generated: 2026-11-03 13:35:19

Machine: DESKTOP-39I5E6Q (Windows 11)

`Project root: C:\Users\HP\Downloads\dr-project-python-fixed\dr-project-python-master`

This document records the work done to take the rebuilt dr/data/ package from an untested checkout to a running, resumable, full-scale training pipeline on this machine - every real bug found along the way, the checkpoint/resume architecture added because this machine cannot stay powered on for the full multi-day run, the CSV export additions, and a live snapshot of where the pipeline stands as of the timestamp above. Re-run this script any time to refresh that snapshot.

1. Environment Setup

The project ships requirements.txt pinned to torch>=2.2 with no Python version constraint; this machine's only available interpreter is Python 3.14, which is new enough that the default PyTorch CUDA wheel index (cu121) has no build for it at all.

- Created .venv on the system Python 3.14 interpreter.
- Installed PyTorch 2.11.0 + torchvision from the cu128 wheel index (download.pytorch.org/whl/cu128) - the cu121/cu124 indices either lack a 3.14 build or lack Blackwell (RTX 50-series) kernel support.
- Installed the remainder of requirements.txt (timm, opencv-python-headless, scikit-learn/-image, pandas, matplotlib, onnx/onnxruntime, kaggle) - all resolved cleanly for 3.14.
- Placed the user-supplied Kaggle API credentials at %USERPROFILE%\.kaggle\kaggle.json so scripts/11_download_lesions.py could fetch the real IDRiD and DDR lesion-annotation corpora.

2. GPU Functionality

Verified end-to-end before committing any real compute time: package installation alone does not prove the CUDA kernels actually match this GPU's architecture.

PyTorch build: 2.11.0+cu128

CUDA available: True

Device: NVIDIA GeForce RTX 5060

VRAM: 8.5 GB

Compute capability: sm_120

Verification performed: allocated a tensor on the CUDA device and ran a real matrix multiply on it (not just an availability check) - confirmed the cu128 build's sm_120 kernels execute correctly on this Blackwell-architecture card, which is not guaranteed by an older CUDA build simply importing without error.

Constraint this GPU imposes on the rest of the run: the project's own full-schedule script (run_gpu_full.sh) assumes >=16GB VRAM at batch-size 32 for a ViT-Large backbone across a 448px global view plus 6 local crops. This card has 8GB, so the training stage uses --batch-size 4 --grad-accum 8 --grad-checkpoint (same effective batch of 32, same total compute, more but smaller steps) as the starting point - see resume_pipeline.ps1.

Observed throughput on this GPU/CPU combination so far:

- Preprocessing (CPU-bound, 24 worker processes): ~1.7-2.0 images/sec sustained on real 1024px EyePACS images.
- Lesion-expert pretraining (GPU-bound, small model): ~8-9 images/sec at 448px on the real IDRiD/DDR crop cache.
- Full staged training is running on the prepared cache; the live training step and throughput are recorded in section 6 and the experiment-values appendix.

3. Real Bugs Found and Fixed

This code had clearly never been run end-to-end against real, full-scale data before - every one of these surfaced only once actual EyePACS/IDRiD/DDR data and a real training step were exercised, not from the project's own synthetic self-test.

3.1 src/dr/data/eyepacs.py

Problem: build_manifest() only searched a few hardcoded subfolders for images and never found this machine's actual layout (resized_train/resized_train/); it silently found 0 of the 35,126 images.

Fix: Added the actual nested layout (and its _cropped variant) to the search list.

3.2 src/dr/data/lesion_datasets.py

Problem: Folder discovery picked the FIRST matching image/mask folder found anywhere in the tree. For IDRiD (three sections, each with its own 'Original Images' folder) this paired images from the wrong section against the mask folder; for DDR (train/test/valid splits) it paired one split's images against another split's masks. Both cases produced 0 (image, mask) pairs.

Fix: Rewrote discovery to merge every matching folder by filename stem instead of picking one arbitrarily. Now correctly finds all 81 IDRiD and 757 DDR annotated images.

3.3 scripts/08_selftest.py

Problem: The crop-centre assertion expected normalised 0-1 coordinates, but generate_lesion_crops() returns pixel coordinates in the cache's own pixel grid (confirmed against how 02_preprocess.py and CachedEyePACS actually consume the values) - a self-test bug, not a production bug.

Fix: Fixed the assertion to check against the cache_size bound instead of [0,1].

3.4 src/dr/modules/a4_backbone_lora.py

Problem: When require_retfound=False and no checkpoint could be resolved at all, load_retfound_weights() returned a partial report dict; the status-string builder unconditionally read fields (matched_tensors, etc.) that only exist on a successful load, crashing with KeyError.

Fix: Branch on whether the load actually succeeded before building the status string; fall back to a clear 'not loaded' message otherwise.

3.5 scripts/10_lesion_pretrain.py

Problem: Called train_ds.lesion_prevalence() as a method; it is a @property already returning a dict, so calling it raised TypeError. The print statement right behind it then zipped a 6-name tuple against the resulting 4-key dict, which would have crashed again on the next attempt. Separately, the training loop read batch keys 'lesion_masks'/'presence' that the dataset never produces (it yields 'masks'/'valid', with no separate presence label at all).

Fix: Removed the erroneous call, fixed the print to iterate the dict directly, and derived presence from the masks (per-channel: any positive pixel in the crop).

3.6 src/dr/modules/external_lesion_batches.py and scripts/05_predict.py

Problem: Both independently treated generate_lesion_crops()'s pixel-space centres as 0-1 fractions and multiplied by image height again, pushing every centre out of range. Since both clip bounds are equal in that regime, every crop for a given image collapsed onto the same fixed corner regardless of where the annotated lesion actually was - this fed both the real-mask auxiliary training path and the Objective 4 XAI evidence path.

Fix: Removed the erroneous re-multiplication in both files.

3.7 scripts/03_train.py

Problem: The training loop read a batch key called 'index' that CachedEyePACS never produces (it's called 'cache_row') - would have crashed on the first training step.

Fix: Changed the read to 'cache_row'.

3.8 scripts/03_train.py

Problem: HardExampleState.empty(len(train_ds.meta)) sized the hard-example-mining arrays to the TRAIN-SPLIT count (~28,102), but they are indexed by cache_row, a global row index assigned before the train/val/test split (so it ranges over the whole ~35,126-row cache). Any sample whose cache_row fell outside the narrower train-count range would index out of bounds.

Fix: Sized the state to len(train_ds.meta_full) instead, matching how CachedEyePACS already sizes its own internal hardness buffer for the same reason.

4. Checkpoint / Resume Architecture

This machine cannot stay powered on for the multi-hour preprocessing pass or the multi-day training run, so neither script could originally survive an interruption - a shutdown mid-run meant restarting from image 1 or epoch 1. Both now checkpoint and auto-resume.

4.1 scripts/02_preprocess.py

- Writes one row to preprocess_progress.csv per successfully processed image (flushed + fsynced every 200 images), and a separate preprocess_failed.json for permanently-unreadable/errored images so they are excluded from the dataset without being retried forever on every resume.
- On startup: if the cache files already exist at the exact shape this invocation implies, opens them in r+ mode (instead of overwriting) and skips every already-completed or already-failed index.
- meta.csv (the file training actually reads) is written to a .tmp path and atomically renamed into place, so a crash mid-write can never leave a corrupt/partial meta.csv.

- Verified live: stopped the real run after ~400 images, restarted the exact same command, and confirmed it printed '400/35126 already cached, continuing with the remainder' instead of reprocessing from scratch.

4.2 scripts/03_train.py

- Saves a full checkpoint.pt after EVERY completed epoch (not only new-best epochs): model weights, optimizer state, LR-scheduler state, EMA shadow weights, the hard-example-mining state, training history, best/best_epoch/stale, and the torch/CUDA/numpy RNG states - written to a .tmp file and atomically renamed, so a shutdown mid-save cannot corrupt the one checkpoint a resume depends on.
- On restart: reconstructs the LoRA adapter architecture deterministically from the saved per-block ranks (bypassing the RNG-sensitive GLA-LoRA calibration pass entirely, since the LoRA layers are real injected modules that model.load_state_dict() needs to already exist), restores every piece of state above, skips any fully-completed stage, and resumes the partially-completed stage at the next epoch with its original optimizer/scheduler state.
- Worst-case loss from an unplanned shutdown: the single epoch in progress, never more.

4.3 scripts/10_lesion_pretrain.py

Given the same per-epoch checkpoint treatment for consistency, though its crop-cache reuse already made a full re-run cheap.

4.4 resume_pipeline.ps1

A single launcher at the project root chaining self-test -> preprocessing -> lesion pretraining -> staged training -> the full evaluation/reporting pipeline. Safe to re-run after any restart - each stage auto-detects and continues from its own checkpoint. Usage:

```
powershell -ExecutionPolicy Bypass -File resume_pipeline.ps1
```

Caveat: once training has written its first checkpoint, the arguments inside this script (particularly --batch-size/--grad-accum) should not change - a resumed run reuses the saved LR-schedule state, which a change would throw off. The final evaluation/report stages are not individually checkpointed (only the two multi-hour stages are); if one of those is interrupted, only that one step needs re-running.

5. CSV Export Additions

A shared helper, src/dr/csv_export.py (save_csv_alongside), writes a best-effort tabular CSV next to whichever JSON a script already produces, wherever the data has one obvious shape: a list of flat dicts becomes one row per entry; a dict of dicts becomes one named row per key; a flat dict becomes a single summary row. Verified against each of these shapes before wiring it into every script below.

- 03_train.py: history.csv (per-epoch curve), gla_lora.csv (per-block LoRA ranks/importance)
- 10_lesion_pretrain.py: history.csv, sources.csv
- 07_compare_models.py: comparison.csv (QWK-sorted leaderboard), results.csv (every field), significance.csv
- 04_evaluate.py: evaluation.csv (long format: group/metric/value), evaluation_confusion_matrix.csv
- 06_deploy.py: distilled_metrics.csv, distilled_variants.csv
- 09_screening.py: screening.csv, screening_deployable.csv

- 12_lesion_visibility.py: lesion_visibility.csv
- 13_linear_probe.py: linear_probe_<tag>.csv
- 14/15/18/20/22_*.py: one CSV per objective evidence-pack JSON
- 16_ordinal_calibration.py: ordinal_calibration_results.csv
- 19_objective3_validation.py: objective3_validation_arms.csv
- 21_xai_evaluation.py: xai_evaluation_per_image.csv, xai_evaluation_metrics.csv
- 23_objectives_summary.py: objectives_consolidated.csv

6. Current Pipeline Status (live at generation time)

6.1 Preprocessing

Images cached so far: 35,126 / 35,126

Last log lines:

```
[preproc] 14901/35126  1.9 img/s (this run)  eta 180.3 min
[preproc] 15401/35126  1.9 img/s (this run)  eta 177.1 min
[preproc] 15901/35126  1.9 img/s (this run)  eta 172.4 min
[preproc] 16401/35126  1.9 img/s (this run)  eta 167.8 min
```

6.2 Lesion-Expert Pretraining

COMPLETE - 8 epochs run, best validation Dice 0.1401. Checkpoint: outputs/lesion_pretrain/lesion_encoder.pt

6.3 Staged Training

IN PROGRESS / RESUMABLE - 31 epoch(s) checkpointed so far.

```
Last completed epoch: 31  stage=s3_unfreeze  QWK=0.5786  score=0.5067
```

```
Latest logged training step: stage 4 epoch 1, step 950/3000, loss=4.147, throughput=2.36 images/sec.
```

6.4 Disk / Environment

Free disk space: 615.5 GB

GPU: NVIDIA GeForce RTX 5060 (8.5 GB VRAM)

7. Next Steps

- Training is currently logged at stage 4 epoch 1, step 950/3000. The latest completed epoch is s3_unfreeze epoch 14 (run epoch 31); refresh this report after the next epoch completes to capture its validation metrics.
- Task Scheduler boot-trigger fix is applied and confirmed working: the machine was off ~2 days (Oct 12 evening - Oct 14), and auto_start_training.ps1 fired on its own at boot with no manual intervention, resuming stage 3 epoch 8 mid-epoch at step 1300/1500 rather than redoing it.
- The full held-out evaluation and post-training evidence pack remain pending until staged training finishes.
- Once training stops (by completing all four stages or an intentional stop), run the full

evaluation/reporting sequence (already chained in resume_pipeline.ps1) and read outputs/objectives_consolidated.json / .csv first.

- Re-run this script (generate_session_report.py) at any point to refresh this document with current progress.

Detailed Methodology – The Logic Behind Every Mechanism

This section exists for one purpose: so the reasoning behind every design choice in this system can be understood, explained, and defended on its own terms - not taken on faith. For each mechanism below: what problem it solves, how it works mechanically (the actual computation, not just the name), and why this particular design was chosen over the obvious simpler alternative. Where real training data already shows the mechanism working, that evidence is given immediately alongside the explanation.

A. Framing: Why DR Grading Is Not Ordinary Classification

Diabetic retinopathy severity is graded 0-4 (No DR, Mild, Moderate, Severe, Proliferative). Treating this as five unrelated classes - the way a generic image classifier would - throws away the most important fact about the problem: the grades are ORDERED, and mistaking grade 0 for grade 4 is a far more serious error than mistaking grade 2 for grade 3. A plain softmax classifier has no concept of this - a confident wrong answer of grade 4 when the truth is grade 0 costs it exactly the same loss as confusing grade 2 and grade 3. Every ordinal-aware component in this system (the CORAL head, the boundary loss, the QWK selection metric) exists because of this single observation, and it is worth holding onto while reading everything that follows.

B. Objective 1 – Class Imbalance: The Full Logic

B.1 The problem, quantified

The real training split (28,102 images) breaks down as 20,641 / 1,944 / 4,250 / 707 / 560 for grades 0-4. Grade 0 outnumbers grade 4 by 36.9 to 1. A model trained on this data with no correction at all has a trivial way to minimise its loss: predict grade 0 for everything. That model would be right 73% of the time (matching the class prior) while being clinically useless - it would never once flag a sight-threatening case. This is not a hypothetical: it is the single most common failure mode reported for DR grading models trained naively, and it is the concrete problem every mechanism in this section exists to prevent.

B.2 Moderate re-sampling: the actual formula

Standard practice for imbalanced classification is inverse-frequency sampling: draw class c with probability proportional to 1/n_c, so each class contributes roughly equally per epoch regardless of how many images it actually has. This project's own prior analysis found that FULL inverse-frequency weighting over-corrects here: with a 36.9x imbalance, 1/n_c weighting makes grade 4 so dominant in the effective gradient that the model swings the other way and starts over-predicting severe disease, while the middle grades (Mild, Moderate) - which are neither the large majority nor the extreme minority - get squeezed out and collapse. The fix implemented here is a softer exponent: sampling weight proportional to n_c^(-0.5) instead of n_c^(-1). Concretely, for this dataset's counts, that means grade 4 gets roughly sqrt(20641/560) = 6.07x the per-image sampling weight of grade 0 - a real, substantial boost, but nowhere near the 36.9x that full inverse-frequency weighting would apply. This is a deliberate middle path: enough correction to force the

model to engage with rare grades, not so much that it forgets grade 0 exists.

- Self-test verification: with this exact formula, the automated self-test computes a grade4/grade0 sampling-weight ratio of exactly 5.01 on its synthetic class counts - confirming the formula is implemented correctly, independent of whether it helps on real data.

B.3 Hard-example mining: what it adds on top

Class-frequency-based re-weighting treats every image of a given grade identically. But within grade 2 (Moderate), some images are textbook-obvious and some sit right on the boundary with grade 1 or grade 3 - those boundary cases carry far more training signal. Hard-example mining tracks THREE per-sample signals across epochs, each updated as an exponential moving average (EMA, decay 0.7, so a single noisy batch cannot spike a sample's priority): (1) the sample's own loss relative to the epoch's mean loss, (2) whether the last prediction on it was a boundary error (off by exactly one grade - the specific No-DR/Mild, Mild/Moderate, Moderate/Severe, Severe/PDR confusions that matter clinically), and (3) whether the top-two predicted class probabilities were close together (a low-margin, low-confidence prediction). A sample's final sampling weight multiplies the class-frequency term by min((1+H)^0.5, 4.0) - where H is the blended hardness score - times a further 1.6x bonus if it was a boundary error and 1.4x if it was low-confidence. The clip at 4.0 stops one pathologically hard (or mislabeled) sample from dominating an entire epoch.

- Self-test verification: the mining mechanism is confirmed to lift the sampling weight of Mild/Moderate boundary-error samples by 1.82x on synthetic data with known hardness scores.

B.4 The CORAL ordinal head: turning severity into K-1 yes/no questions

Instead of a 5-way softmax, the grading head is CORAL (Consistent Rank Logits): it asks K-1=4 binary questions - 'is severity > 0?', '> 1?', '> 2?', '> 3?' - each sharing the same underlying feature representation but with its own learned threshold (a 'ladder' of 4 cutoffs, b_0 through b_3). The predicted grade is simply how many of these thresholds the sample's latent severity score exceeds. This guarantees rank-consistent predictions (the probability of 'severity > 1' can never exceed the probability of 'severity > 0', which a naive independent multi-task setup would not guarantee) and, more importantly for this objective, it means the DECISION BOUNDARY for each grade transition is a SEPARATE, individually-adjustable parameter - exactly the lever needed to correct for a dataset where those transitions occur at very different frequencies.

B.5 Why the ladder is seeded from the real class prior (the '1e-4' finding)

A CORAL ladder is normally initialised with a uniform gap between thresholds (evenly spaced, data-independent). This project's own evidence pack traced a real, previously undiagnosed failure to exactly this default: under uniform initialisation, the learned thresholds were found to move by only about 1e-4 over an ENTIRE training run - three orders of magnitude too small a change to meaningfully reshape which images fall into which grade. The practical consequence was that Mild and Severe recall pinned at zero no matter how the sampler or loss weights were tuned, because the decision boundaries were never actually moving to where the minority-grade images live in score-space. The fix: initialise each threshold directly from the empirical class prior, b_k = logit(P(Y>k)), computed from the real training counts. On THIS dataset's counts, that means the ladder starts at [0.0, -0.974, -1.948, -2.922] under the old uniform scheme versus [-1.018, -1.410, -3.053, -3.895] under prior-initialisation - already spread across a realistic ~2.9-logit range before a single gradient step, instead of needing training to discover that spread from a near-zero start. The ladder is additionally trained at 10x the head's learning rate, attacking the same root cause (thresholds that barely move) from the other direction as well.

B.6 DQK loss: optimising the actual evaluation metric, not a proxy

Quadratic Weighted Kappa (QWK) is the primary metric this project is judged on, but QWK itself is not differentiable (it is a discrete statistic over a confusion matrix), so it cannot be used directly as a loss function. DQK (Dice-weighted ordinal loss) is a differentiable surrogate: it builds a SOFT confusion matrix from the model's predicted probabilities (rather than hard argmax predictions), weights each cell by the same quadratic distance-from-diagonal penalty QWK itself uses, and minimises that directly. This closes the gap between 'what the loss function optimises' and 'what the model is actually evaluated on' - standard cross-entropy-style losses optimise per-sample correctness, not the AGGREGATE ordinal-agreement statistic the project is scored on.

B.7 How all four mechanisms combine, and why that is not double-counting

Sampling re-weighting, hard-example mining, prior-initialised CORAL, and DQK loss operate at three DIFFERENT stages of the pipeline - which data gets shown to the model (sampling+mining), how the model's own decision boundaries start out (CORAL prior-init), and what the model is pushed to optimise once it sees a batch (DQK plus the other loss terms). None of them alone would be sufficient: re-weighting alone does nothing about badly-placed decision thresholds; a well-placed threshold alone does nothing if the model never sees enough minority examples to learn what distinguishes them; and neither helps if the loss function being optimised is misaligned with the evaluation metric. They are layered, not redundant.

B.8 The evidence, read as a narrative

Reading the five real stage-1 epochs in order: epoch 1 recall per grade was [0.829, 0.059, 0.224, 0.278, 0.000] - grade 4 (the rarest, most severe class) was NEVER once correctly identified, exactly the failure mode this whole section is built against. By epoch 5, recall was [0.758, 0.059, 0.379, 0.194, 0.429] - grade 4 recall rose from 0% to 42.9%, grade 2 rose from 22.4% to 37.9%, while grade 0 recall FELL from 82.9% to 75.8%. That fall is not a defect - it is the imbalance-handling machinery visibly doing its job: the model is trading a little accuracy on the class it could get right for free in exchange for real traction on the classes that actually matter clinically. The one class that has not moved yet is grade 1 (Mild), still at 5.9-10.9% recall throughout - a known hard case in DR grading generally (visually intermediate between No-DR and Moderate), and the clearest open question as training continues into stages 2-4, where more of the network becomes trainable.

C. Objective 2 - Preprocessing: The Full Logic

C.1 Why one fixed preprocessing recipe fails

A single global contrast/illumination correction, applied identically to every image, is fighting two sources of variation at once: camera/lighting variation (some images are already well-exposed, some are dark or glare-heavy) and pathology variation (a barely-visible microaneurysm needs different handling than an obvious haemorrhage). A fixed recipe strong enough to rescue the worst images over-processes the good ones (introducing artefacts, blowing out contrast); a recipe gentle enough to leave good images alone under-corrects the bad ones. The system instead makes TWO separate decisions per image: how much correction does this image need (A1), and which blend of correction techniques suits it best (A2/ALPP) - both computed FROM the image itself, not fixed in advance.

C.2 The five Q* quality axes - what each one actually measures

- Q_quality: a general well-formedness score combining gradability, artefact presence, and field-of-

view completeness (sub-weights 0.5/0.28/0.22) - is this even a usable fundus photograph at all.

- Q_domain: how far this image's colour/illumination statistics (red-green channel ratio, saturation, fill fraction) sit from the calibrated reference for this camera/dataset - a proxy for 'does this look like a normal image from this acquisition source, or an outlier'.

- Q_lesion: contrast-to-noise ratio of lesion-sized structures against local background, calibrated against a measured floor and reference CNR for this corpus - directly, this is 'can lesions actually be seen in this image', which is why it is expected to rise with DR grade (more severe disease -> more, more visible lesion evidence) and, on the real data, mostly does: 0.453 at grade 0 rising to 0.559 at grade 4.

- Q_blur: a sharpness measure (relates to high-frequency energy in the image, calibrated against a reference blur level for this corpus).

- Q_illumination: how evenly and adequately the retina is lit across the field, independent of the domain/colour check above.

These five combine (weights 0.3/0.1/0.25/0.2/0.15 for quality/domain/lesion/blur/illumination) into one Q* score, which is compared against two thresholds (tau_r=0.32 for outright rejection, tau_d=0.6 for 'good enough to accept without enhancement') to route each image to accept / enhance / retake. On the real 35,126-image dataset: 52.7% needed no correction at all, 46.9% were routed to enhancement, and only 0.43% were flagged as unrecoverable - meaning the enhancement path (A2/ALPP below) is doing real, substantial work on nearly half the dataset, not a rare edge case.

C.3 ALPP: four correction branches, one learned gate

Rather than picking ONE enhancement algorithm, four candidate corrections are computed for every enhancement-routed image: illumination normalisation (norm), adaptive CLAHE with a clip limit of 2.5 on an 8x8 grid (contrast), a vessel-enhancement pass, and a lesion-targeted enhancement pass. A small gate network then decides, PER IMAGE, how to blend these four - not a fixed recipe. The gate works like this: the image is downsampled to 128x128 (verified directly in the source: F.interpolate(..., size=(128,128))) so the gate only has to reason about coarse, global structure - illumination pattern, general contrast - not pixel-level lesion detail, which keeps the gate small (64 hidden units) and fast. That 128x128 summary, combined with the image's own 5-axis quality scores, is fed through the small network to produce four gate weights that sum to 1 (a softmax), which is why 'Sum(g)=1' is a checkable invariant, not just documentation.

The gate's initial bias is set to (1.2, 0.6, 0.3, 0.3) BEFORE any training - a deliberate prior favouring illumination-normalisation and CLAHE over the vessel/lesion branches at initialisation, since those two are the safest, most broadly-applicable corrections; the network then learns to deviate from this prior per image as it trains. A final fusion parameter gamma=0.6 controls how strongly the blended correction is applied versus the original pixels (a partial-strength blend rather than a full replacement, so the correction cannot introduce more artefacts than it removes).

- Self-test verification: feeding the gate a synthetic dark image and a synthetic bright image produces MEASURABLY DIFFERENT branch weights for each (confirmed, not assumed) - proving the gate is actually conditioning on the image rather than converging to one constant blend regardless of input.

C.4 Native-resolution cropping: the actual pixel arithmetic

A microaneurysm is tens of micrometres across - a genuinely tiny structure. If a whole fundus image is simply resized down to a practical training resolution (say 224px) before any lesion detection happens, that resize operation itself can shrink a microaneurysm to a fraction of a single output pixel, at which point no

model, regardless of architecture, can learn to detect it - the information was destroyed by the resize, not by the model's limitations. This system instead cuts a smaller window (320px, in the cache's own native 1024px pixel grid) CENTRED on lesion-dense regions (found via the weak morphological priors computed at cache time), and only resizes that already-small window down to the network's input size (224px). The lesion is sampled at 224/320 = 0.70 output pixels per cache pixel this way, versus 224/1024 = 0.22 for a whole-image view - roughly 3.2x more effective resolution over exactly the region that matters.

D. Objective 3 – Backbone and Efficient Fine-Tuning: The Full Log

D.1 Why a retina-specific foundation model, not ImageNet

A backbone pretrained on ImageNet (photographs of everyday objects) has learned features tuned for natural-image statistics - edges, textures, and shapes that distinguish a dog from a car. Fundus photographs have almost none of that structure in common: the relevant signal is subtle colour and contrast variation against a highly repetitive circular background. RETFound is a Vision Transformer pretrained with masked-autoencoding (MAE) directly on large collections of retinal images - the pretraining TASK is 'reconstruct randomly masked patches of a retinal image from the visible patches', which forces the network to learn what normal retinal structure looks like, so it already has a head start on noticing when something is abnormal, before a single labelled DR image is shown to it. This is why the coverage check matters so much: 100% of the 294 backbone weight tensors loaded correctly, confirming the ACTUAL published RETFound MAE weights are in use, not an ImageNet fallback silently substituted in.

D.2 LoRA in one paragraph

Fully fine-tuning a 319.6-million-parameter model risks catastrophic forgetting (the retina-specific knowledge from pretraining gets overwritten before the DR-grading task has a chance to benefit from it) and is expensive per step. LoRA (Low-Rank Adaptation) instead freezes the original weight matrix W entirely and learns a small correction on the side: W_effective = W + (alpha/r) * B*A, where A and B are small matrices of rank r (here, r ranges 2-16 depending on the block - see below) - so the number of NEW trainable parameters per layer is proportional to r, not to the full size of W. This gives most of the benefit of fine-tuning at a small fraction of the parameter cost, and W itself never changes, so the pretrained knowledge is preserved by construction, not by hoping the learning rate stays small enough.

D.3 GLA-LoRA: why every block gets a DIFFERENT rank

Plain LoRA uses the same rank for every transformer block, which implicitly assumes every block matters equally for the task - an assumption with no real justification. GLA-LoRA (Gradient/Lesion/Attention-guided LoRA) instead computes THREE separate importance signals for each of the 24 blocks, on a real calibration batch, BEFORE training starts:

- Gradient-sensitivity (G_l): how much the loss changes with respect to this block's output - a large gradient means small changes to this block's behaviour would meaningfully move the loss, so it is a block worth being able to adapt.
- Lesion-relevance (L_l): how much this block's activations respond specifically to lesion-containing regions of the input, versus responding uniformly regardless of content.
- Attention-mass (A_l): what fraction of this block's own attention weight concentrates onto lesion regions rather than spreading over background retina.

These three are each min-max normalised to [0,1] and blended: S_l = 0.35*G_l + 0.45*L_l + 0.20*A_l (lesion-

relevance weighted highest, since this is a lesion-driven diagnostic task). The rank for each block is then r_l = r_min + (r_max - r_min) * S_l, clamped to [2, 16], with an extra rule: any block whose S_l exceeds 0.85 is pinned straight to the maximum rank regardless of the interpolated value (the 'lesion-sensitive pin'), so a block that is unambiguously important never gets shortchanged by interpolation. On THIS run's real calibration batch: block 0 (highest gradient-sensitivity, G_l=1.0) got rank 13; block 23 (the final block, S_l collapses to 0.094 with ZERO lesion-relevance) got the floor rank of 3. The allocation is genuinely non-uniform and importance-driven - not a formality.

Total cost: 1,431,552 trainable LoRA parameters against the 319.6M-parameter backbone - 0.45%. That is the concrete, measured form of 'lightweight optimisation without full retraining'.

D.4 The 4-stage schedule: escalating only as far as needed

Stage 1 (5 epochs) trains only the head, with the entire backbone frozen - the cheapest possible adaptation, establishing a baseline before anything in the pretrained network is touched. Stage 2 (12 epochs) activates the LoRA adapters (and lets ALPP start training too) while the backbone itself stays frozen - still cheap, but now the network can make small, low-rank adjustments to its internal representations. Stage 3 (14 epochs) unfreezes the LAST 33% of the transformer blocks (the ones closest to the output, which encode the most task-specific features) at a small learning rate (5e-6) alongside the LoRA adapters. Stage 4 (14 epochs) unfreezes the entire backbone at the same conservative backbone learning rate. Each stage only escalates to the next, more expensive, more forgetting-prone level of adaptation once the cheaper level has had its full chance - the schedule is the literal mechanism behind 'efficient adaptation... without full retraining': full retraining is the LAST resort here, not the starting point.

D.5 Lesion-expert pretraining and transfer

The lesion-evidence experts (one per lesion channel: microaneurysms, haemorrhages, hard exudates, soft exudates, plus two channels - neovascularisation and macular oedema - that neither IDRiD nor DDR annotate and are therefore masked out of every loss involving them) are pretrained SEPARATELY, on real pixel-level annotations from 838 IDRiD/DDR images, at the SAME magnification the main model's local crops will use (31.2% of the retinal field per window) - matching magnification matters because a lesion expert trained on whole, downsampled images never sees a microaneurysm at a resolvable size, so its Dice score pins at zero regardless of how long it trains. Only the convolutional stem and the six expert heads transfer into the main model (124 tensors, 0 missing, 0 unexpected); the adaptive gate does NOT transfer, because it conditions on signals (image quality, hard-example history) that only exist once real DR-grading training is underway, not during lesion-only pretraining.

E. Objective 4 - Explainability, Trust, and Deployment: The Full L

E.1 Attribution during training, not only at evaluation

A model can be scored for explainability after the fact (does its saliency map overlap real lesion masks) without that requirement ever influencing training - which means a model can score well on accuracy while attributing its decisions to irrelevant regions (a known failure mode called 'right answer, wrong reason'). This system instead adds an attribution-consistency loss DURING training: the model's own Grad-CAM-style attribution map is compared against the lesion evidence map and penalised for disagreement, with a weight lambda_XAI. The default schedule holds lambda_XAI at 0 through stages 1-2 (attention is too noisy early in training to usefully supervise) and only turns it on at stage 3 - but the project's own evidence pack found a run that NEVER REACHED stage 3 therefore trained with lambda_XAI=0 for its entire duration, and scored at

chance on attribution as a direct, mechanical consequence, not a deeper flaw. This run overrides lambda_XAI to a constant 0.05 from stage 1 specifically so that exact failure cannot repeat.

E.2 The real-mask fix (Objective 4's own documented correction)

EyePACS itself carries no pixel-level lesion annotations, so attribution can only be scored against WEAK morphological priors on it (a top-hat-filter-style heuristic, not ground truth). This system instead scores and supervises attribution against the same 838 REAL ophthalmologist-annotated IDRiD/DDR images already used for lesion-expert pretraining, folded in as an auxiliary batch every 5 training steps. This replaces a weak-reference caveat with a real one - the numbers this produces are directly comparable to published lesion-attribution results elsewhere, not qualified by 'against a proxy label'.

E.3 The counterfactual test: a causal check, not a correlational one

A high overlap between an attribution map and a lesion mask shows correlation - the model looked at the right place - but not that the model's DECISION actually depended on what it saw there. The counterfactual test closes this gap: the cited lesion region is erased (inpainted) from the image, and the model is asked to re-predict severity. If the prediction does not meaningfully change, the original attribution was not causally load-bearing, regardless of how well it overlapped the lesion mask visually.

E.4 Uncertainty: MC-dropout, calibration, and the decision gate

Dropout is normally switched OFF at inference for a deterministic prediction. MC-dropout instead keeps it ON and runs the same image through the network multiple times (10 samples here); because dropout randomly zeroes different neurons each pass, the resulting spread of predictions is a direct estimate of the model's own epistemic uncertainty - how much its answer would change under small perturbations to its own computation. Separately, temperature scaling (a single learned scalar dividing the logits before softmax) and isotonic regression are both available to CALIBRATE the model's raw confidence scores against its actual empirical accuracy, since an uncalibrated network's '90% confident' frequently does not mean '90% likely to be correct'. Both signals feed a three-way decision gate - AI_DECISION (high confidence, proceed), HUMAN_REVIEW (uncertain, flag for a clinician), RETAKE_IMAGE (the underlying image quality itself is the problem) - which is the concrete trust-building behaviour this objective calls for: a system that visibly knows when it does not know, instead of forcing every case through the same confident-looking output.

E.5 Deployment: distillation and quantisation

The full model (RETFound ViT-Large plus every module above) is too large for real-time point-of-care deployment. A smaller, image-only student network (no anatomy maps, no lesion MoE, so it can run without the full pipeline) is trained via knowledge distillation - learning from the teacher's soft ordinal probability distribution plus the real labels - then further compressed via FP16/INT8 quantisation and exported to ONNX. Every size and latency number this produces is measured ON THIS machine, not quoted from a paper, specifically so the deployment claim is checkable against real hardware.

Everything in this section is either a mechanism already verified correct in isolation by this project's own automated self-test, or a mechanism whose effect is already visible in the real per-epoch training numbers reported elsewhere in this document - none of it is aspirational or unverified design intent.

Appendix: Experiment Values Registry

Source: data\experiment_values.json, generated 2026-11-03T13:34:34. Every value below is read from a real config/log/output file at generation time; anything not yet computable is marked PENDING, and anything this

codebase does not implement is marked NOT APPLICABLE, rather than estimated or invented.

1 dataset data integrity

dataset name: EyePACS (Kaggle 'Diabetic Retinopathy Detection', high-res 'eyepacs-hi' export) for training/val/test; APTOS 2019 as an external-only evaluation cohort; IDRiD + DDR for lesion annotation supervision (see section 8)

dataset version: Kaggle competition export as downloaded to data/raw/eyepacs-hi (trainLabels.csv, 35,127 labelled rows; 35,126 image files actually present and used)

total image count: 35126

total patient count: 17563

total eye count: NOT APPLICABLE - not implemented in this codebase as a separate figure (each row IS one eye; patient_id is the shared numeric prefix of '<id>_left'/'<id>_right')

image ids patient ids eye ids: stored per-row in data/cache/meta.csv (image_id, patient_id, cache_row columns) - not reproduced here row-by-row (35,126 rows)

original image paths: data/raw/eyepacs-hi/resized_train/resized_train/<image_id>.jpeg

image dimensions native: as downloaded from Kaggle (varies by source camera); resampled into the cache at cache_size below

cache resolution px: 1024

dr grade counts:

0: 25810

1: 2443

2: 5292

3: 873

4: 708

class percentages:

0: 73.48

1: 6.95

2: 15.07

3: 2.49

4: 2.02

train count: 28102

val count full: 3512

val count used for epoch monitoring: 1500

test count: 3512

train class counts: [20641, 1944, 4250, 707, 560]

validation class counts: PENDING - computed after training/evaluation completes (not logged per-class during training; computable from data/cache/meta.csv split=val rows)

test class counts: PENDING - computed after training/evaluation completes (only read at final evaluation, scripts/04_evaluate.py)

split level: patient-level (both eyes of one patient always share a split - src/dr/data/eyepacs.py:split_by_patient)

eye level split: NOT APPLICABLE - not implemented in this codebase (this project deliberately does NOT split by eye, to prevent leakage - see split_level above)

split random seed: 1337

split csvs: NOT APPLICABLE - not implemented in this codebase (the split is a deterministic function of meta.csv + seed, recomputed each run rather than exported as a standalone CSV; the 'split' column IS a permanent field of data/cache/meta.csv)

train val overlap: 0 (patient-level split verified)

train test overlap: 0 (patient-level split verified)

val test overlap: 0 (patient-level split verified)

duplicate image count: NOT APPLICABLE - not implemented in this codebase (no perceptual/hash dedup pass implemented)

duplicate hash count: NOT APPLICABLE - not implemented in this codebase

leakage check result: PASSED - '[split] verified: no patient appears in two splits' (assertion in scripts/02_preprocess.py)

2 preprocessing quality pipeline

global image size px: 448

high res cache size px: 1024

local crop size native px: 320

local crop fed size px: 224

n local crops: 6

crop jitter: 0.12

quality score mean qstar: 0.612

quality score mean legacy: 0.681

q axes mean:

q quality: 0.789

q domain: 0.667

q lesion: 0.453

q blur: 0.436

q illumination: 0.72

q lesion by grade:

0: 0.4531

1: 0.4065

2: 0.4645

3: 0.4417

4: 0.5585

quality by dataset: NOT APPLICABLE - not implemented in this codebase (single-dataset preprocessing run; no per-dataset quality breakdown computed)

gate decisions:

accept:

count: 18501

pct: 52.67

enhance:

count: 16474

pct: 46.9

gate decision per image: stored per-row in data/cache/meta.csv (quality_decision column)

preprocessing runtime per image: ~12.5 CPU-seconds/image single-threaded (measured this session at 1024px, full A1+A2+lesion-priors pipeline)

total preprocessing runtime: ~4.5 hours wall-clock for all 35,126 images with 28 parallel worker processes (measured)

failed preprocessing images: 0 (see data/cache/preprocess_failed.json)

clahe clip limit: 2.5

clahe grid size: 8

illumination norm sigma: 25.0

enhancement algorithm: clahe

large lesion protection:

enabled: True

kernel: 41

strength: 0.6

alpp gate hidden dim: 64

alpp branches: ['norm', 'contrast', 'vessel', 'lesion']

alpp initial bias prior: [1.2, 0.6, 0.3, 0.3]

alpp fusion gamma: 0.6

alpp weights per image by grade by quality: PENDING - computed after training/evaluation completes (per-image ALPP weights are computed on the fly during training/preprocessing, not persisted per-image)

3 alpp verified constants

downsample before gate net px: 128x128 (verified in src/dr/modules/a2_preprocess.py: F.interpolate(..., size=(128,128)))

gate hidden units: 64

initial bias prior: [1.2, 0.6, 0.3, 0.3]

fusion gamma: 0.6

branches: ['norm', 'contrast', 'vessel', 'lesion']

note: these four values were independently cross-checked against the live config.json AND the

a2_preprocess.py source for this exact run and all matched what had been previously documented (128x128 / 64 / [1.2,0.6,0.3,0.3] / 0.6)

4 retfound backbone

checkpoint source: bitfount/RETFound_MAE:pytorch_model.bin

checkpoint fallback configured: YukunZhou/RETFound_mae_natureCFP:RETFound_mae_natureCFP.pth

require retfound: True

load report line: [A4] RETFound weights loaded from C:\Users\HP\.cache\huggingface\hub\models--bitfount--RETFound_MAE\snapshots\0f81b6df4222edacc026f9fab2aa81bc71ebb1cf\pytorch_model.bin (matched 294, missing 0, unexpected 0, backbone coverage 100.0%)

backbone coverage pct: 100.0

architecture: vit_large_patch16_224

input resolution global px: 448

input resolution local crop px: 224

patch size: 16 (vit_large_patch16_224 - patch size encoded in model name)

total model parameters: 319612795

checkpoint hash version: HF snapshot 0f81b6df4222edacc026f9fab2aa81bc71ebb1cf (bitfount/RETFound_MAE)

5 dynamic gla lora

enabled: True

rank min: 2

rank max: 16

lora alpha: 16.0

lora dropout: 0.05

lora targets: ['qkv', 'proj']

importance weights:

lam grad: 0.35

lam lesion: 0.45

lam attn: 0.2

lesion layer boost threshold: 0.85

per block table:

block | importance_S | grad_G | lesion_L | attn_A | rank

0 | 0.7718 | 1.0 | 0.5896 | 0.7824 | 13

1 | 0.4332 | 0.1236 | 0.7277 | 0.3121 | 8

2 | 0.4265 | 0.1795 | 0.8082 | 0.0 | 8

3 | 0.5797 | 0.1029 | 0.8289 | 0.8536 | 10

4 | 0.6459 | 0.164 | 0.8925 | 0.9345 | 11

5 | 0.6313 | 0.0993 | 0.9481 | 0.8493 | 11

6 | 0.6211 | 0.0926 | 0.9073 | 0.9018 | 11

7 | 0.6477 | 0.0561 | 0.9512 | 1.0 | 11

8 | 0.5821 | 0.1103 | 0.9994 | 0.469 | 10

9 | 0.5974 | 0.1663 | 1.0 | 0.446 | 10

10 | 0.6575 | 0.0946 | 0.9861 | 0.903 | 11

11 | 0.5723 | 0.0944 | 0.9645 | 0.5261 | 10

12 | 0.5831 | 0.046 | 0.9674 | 0.6584 | 10

13 | 0.5583 | 0.0401 | 0.9524 | 0.5782 | 10

14 | 0.5351 | 0.0062 | 0.9519 | 0.5229 | 9

15 | 0.4875 | 0.003 | 0.9478 | 0.2995 | 9

16 | 0.614 | 0.0152 | 0.945 | 0.9171 | 11

17 | 0.5576 | 0.0173 | 0.9233 | 0.6802 | 10

18 | 0.4938 | 0.0041 | 0.886 | 0.4684 | 9

19 | 0.5119 | 0.0 | 0.8463 | 0.6553 | 9

20 | 0.5361 | 0.0018 | 0.8584 | 0.7459 | 10

21 | 0.5872 | 0.0169 | 0.8496 | 0.9951 | 10

22 | 0.4704 | 0.0285 | 0.8196 | 0.458 | 9

23 | 0.0941 | 0.009 | 0.0 | 0.4548 | 3

trainable lora parameters: 1431552

backbone total parameters: 319612795

lora parameter pct: 0.4479

n blocks at r max: 0

n blocks at r min: 0

calibration seed: 1337

calibration batch size: --calib-batch (default 4)

documented headline values:

backbone params approx: 319.6M

lora params approx: 1431552

lora pct of backbone: 0.45%

block 0 rank: 13

block 21 rank: 10

block 23 rank: 3

block 23 importance: 0.0941

6 training stages

n stages: 4

stages:

name | epochs | train_backbone | train_lora | unfreeze_frac | head_lr | lora_lr | backbone_lr | train_alpp

s1_frozen | 5 | False | False | 0.0 | 0.0001 | 0.0 | 0.0 | False

s2_lora | 12 | False | True | 0.0 | 0.0001 | 5e-05 | 0.0 | True

s3_unfreeze | 14 | True | True | 0.33 | 5e-05 | 5e-05 | 5e-06 | True

s4_full | 14 | True | True | 1.0 | 2e-05 | 2e-05 | 5e-06 | True

batch size: 2

grad accum: 16

effective batch size: 32

weight decay: 0.02

warmup frac: 0.1

grad clip norm: 1.0

scheduler: warmup + cosine decay (LambdaLR, per-stage step budget)

amp: True

amp dtype: bfloat16 (CUDA Ampere+/RTX 5060 supports bf16 natively)

gradient checkpointing: True

ema enabled: True

ema decay: 0.995

seed: 1337

num seeds: 1

ladder lr mult: 10.0

init ladder from prior: True

num workers: 16

patience epochs: 45

select objective: ordinal

select weights:

qwk: 0.5

macro f1: 0.3

minority recall: 0.2

cpu utilization monitoring: NOT APPLICABLE - not implemented in this codebase per user instruction (not tracked going forward)

data loading forward backward optimizer time breakdown: NOT APPLICABLE - not implemented in this codebase (not separately instrumented; only aggregate throughput img/s is measured)

7 class imbalance sampling

sampling power: -0.5

formula: P_i proportional to n_class(i)^power * min((1+H_i)^hard_eta, hard_clip) * (1+boundary_bonus*boundary_i) * (1+lowconf_bonus*lowconf_i)

hard eta: 0.5

hard eps: 0.001

hard ema decay: 0.7

hard clip: 4.0

boundary bonus: 0.6

lowconf bonus: 0.4

class weight power cb focal: -0.25

class counts train: [20641, 1944, 4250, 707, 560]

self test verified ratio: grade4/grade0 sampling-weight ratio = 5.01 exactly (automated self-test, synthetic counts)

self test verified mining lift: 1.82x boundary-error weight lift (automated self-test)

samples per epoch: 6000 (resume_pipeline.ps1 default)

sampler state checkpointed: no - the sampler/hardness EMA state IS checkpointed (HardExampleState in checkpoint.pt), but the WeightedRandomSampler draw sequence itself is not seeded for exact replay across a resume

8 lesion encoder

idrid images: 81

idrid per channel:

MA: 81

HE: 80

EX H: 81

EX S: 40

ddr images: 757

ddr per channel:

MA: 757

HE: 757

EX H: 757

EX S: 757

total annotated images: 838

nv annotation availability: none (neither IDRiD nor DDR annotate neovascularisation - masked out of every loss)

me annotation availability: none (neither corpus annotates macular oedema)

architecture: LesionMoE (src/dr/modules/a3_lesion_moe.py), 1,824,498 trainable parameters, evidence maps at 112px

pretraining seed: 1337

pretraining epochs: 8

history per epoch:

```
epoch | train_loss | epoch_sec | dice_MA | dice_HE | dice_EX_H | dice_EX_S | dice_NV | dice_ME |
dice_mean | _annotated_per_channel
1 | 30.707891882693723 | 95.9 | 0.0 | 0.0 | 0.13920166771877798 | 0.0 | nan | nan |
0.034800416929694496 | [756, 756, 756, 756, 0, 0]
```

2 | 2.50976720407661 | 80.9 | 0.030419715055833654 | 0.0009455226967658058 | 0.33002379380364505 | 0.0 | nan | nan | 0.09034725788906113 | [756, 756, 756, 756, 0, 0]

3 | 2.240821138303378 | 80.7 | 0.0330642060745867 | 0.022294240654497588 | 0.19218370001696736 | 0.0 | nan | nan | 0.06188553668651291 | [756, 756, 756, 756, 0, 0]

4 | 2.094665002901009 | 82.8 | 0.08115838398283876 | 0.02094127266204457 | 0.31449820811550266 | 0.0 | nan | nan | 0.10414946619009649 | [756, 756, 756, 756, 0, 0]

5 | 2.0033256150363536 | 82.3 | 0.06389193136181089 | 0.055255372729273125 | 0.41991267576580604 | 0.0 | nan | nan | 0.1347649949642225 | [756, 756, 756, 756, 0, 0]

6 | 1.9438250305947293 | 82.4 | 0.09528985507246376 | 0.044079672363032285 | 0.31747131747131746 | 0.0018192844147968466 | nan | nan | 0.1146650323304026 | [756, 756, 756, 756, 0, 0]

7 | 1.904215879598807 | 82.8 | 0.08621626516935338 | 0.04618516755052785 | 0.4210731082148047 | 0.0 | nan | nan | 0.1383686352336715 | [756, 756, 756, 756, 0, 0]

8 | 1.888188412517644 | 82.2 | 0.09419770773638969 | 0.04400411099691675 | 0.42226241443256485 | 0.0 | nan | nan | 0.14011605829146784 | [756, 756, 756, 756, 0, 0]

best epoch: 8

best mean val dice: 0.14011605829146784

final epoch per channel dice:

dice MA: 0.09419770773638969

dice HE: 0.04400411099691675

dice EX H: 0.42226241443256485

dice EX S: 0.0

dice NV: nan

dice ME: nan

dice mean: 0.14011605829146784

transfer into main model: [A3] lesion encoder transferred from outputs\lesion_pretrain\lesion_encoder.pt (124 tensors; missing 0, unexpected 0; pretrained on {'idrid': {'images': 81, 'masks_per_channel': {'MA': 81, 'HE': 80, 'EX_H': 81, 'EX_S': 40}}, 'ddr': {'images': 757, 'masks_per_channel': {'MA': 757, 'HE': 757, 'EX_H': 757, 'EX_S': 757}}}, val Dice 0.1401)

documented headline values:

tensors transferred: 124

missing: 0

unexpected: 0

best mean val dice approx: 0.1401

EX H dice approx: 0.42

MA dice approx: 0.09

HE dice approx: 0.04

EX S dice approx: 0.002

9 cross attention pathology anatomy a5

note: this codebase's cross-attention module is A5 (pathology<->anatomy fusion, src/dr/modules/a5_fusion_gnn.py) - bidirectional attention between lesion-evidence tokens and named

anatomical regions, NOT a separate 'global/local/HR' three-way cross-attention (that describes a different architecture this codebase does not have)

n heads: 8

anatomy regions: ['optic_disc', 'macula', 'sup_arcade', 'inf_arcade', 'peripheral']

gnn layers: 2

knn: 6

fused dim: 384

bidirectional: True

probe dim: 128

self test verified shapes: P->A (2,8,53,5) A->P (2,8,5,53) L_PA=1.7229 (automated self-test, synthetic batch of 2)

global lesion hr specific attention maps: NOT APPLICABLE - not implemented in this codebase

10 ordinal coral

n classes K: 5

n thresholds: 4

threshold lr multiplier: 10.0

init from prior: True

uniform init thresholds: [0.0, -0.9740999937057495, -1.948199987411499, -2.9221999645233154]

prior matched thresholds: [-1.0176000595092773, -1.409500002861023, -3.053100109100342, -3.8954999446868896]

decode rule: coral (rank-counting; --decode argmax also available as an ablation)

pi coral vs standard coral: this run uses prior-initialised CORAL ("PI-CORAL" in the objective-1 write-up) by default (init_ladder_from_prior=true); standard uniform-init CORAL is available via --no-init-ladder-from-prior

reachability bounds per grade: PENDING - computed after training/evaluation completes (computed by scripts/16_ordinal_calibration.py and scripts/15_objective1_report.py's decode-reachability sweep, both post-training)

11 dqk loss

enabled: True

weight w dqk: 0.5

dqk vs direct qwk agreement: PENDING - computed after training/evaluation completes (this is exactly what scripts/15_objective1_report.py's decode-ablation compares, post-training)

12 other loss components

learned loss weighting: True

fixed weights if not learned:

- w ordinal: 1.0
- w hierarchical: 0.5
- w boundary: 0.2

w contrastive: 0.2

w lesion: 0.3

w pa: 0.1

w cbf: 0.3

w dqk: 0.5

focal gamma: 2.0

cb beta: 0.999

boundary gamma: 1.5

contrastive temperature: 0.1

lambda xai schedule: [0.05, 0.05, 0.05, 0.05]

lambda xai override this run: constant 0.05 from stage 1 (--xai-weight 0.05, overriding the default 0-at-stage-1/2 schedule) - deliberate, to avoid the documented lambda_XAI=0 failure mode

emd margin losses: NOT APPLICABLE - not implemented in this codebase (this codebase does not implement separate EMD or margin loss terms - boundary_loss/contrastive_loss serve an analogous role)

13 main training history so far

epochs completed this run:

stage | epoch | train_loss | qwk | f1_macro | min_recall | score | accuracy | ref_auc

s1_frozen | 1 | 83.8989 | 0.3505 | 0.292 | 0.1872 | 0.3003 | 0.666 | 0.726

s1_frozen | 2 | 48.4253 | 0.4205 | 0.2865 | 0.1983 | 0.3359 | 0.6073 | 0.7778

s1_frozen | 3 | 33.6463 | 0.4408 | 0.2842 | 0.1676 | 0.3392 | 0.634 | 0.7935

s1_frozen | 4 | 27.291 | 0.4033 | 0.3012 | 0.2891 | 0.3498 | 0.5687 | 0.7996

s1_frozen | 5 | 25.5096 | 0.4576 | 0.309 | 0.2108 | 0.3636 | 0.6387 | 0.8078

s2_lora | 1 | 24.0468 | 0.4338 | 0.2866 | 0.2283 | 0.3486 | 0.5387 | 0.8089

s2_lora | 2 | 17.4669 | 0.4957 | 0.3757 | 0.2688 | 0.4143 | 0.7247 | 0.8041

s2_lora | 3 | 12.1453 | 0.4651 | 0.3027 | 0.2378 | 0.3709 | 0.6093 | 0.8126

s2_lora | 4 | 9.3385 | 0.4558 | 0.3182 | 0.3014 | 0.3836 | 0.5873 | 0.8204

s2_lora | 5 | 8.1316 | 0.5542 | 0.3989 | 0.3695 | 0.4707 | 0.7033 | 0.8317

s2_lora | 6 | 7.3376 | 0.5536 | 0.39 | 0.3339 | 0.4606 | 0.6907 | 0.8364

s2_lora | 7 | 6.9499 | 0.4023 | 0.3437 | 0.474 | 0.3991 | 0.4793 | 0.836

s2_lora | 8 | 6.6168 | 0.5088 | 0.3797 | 0.3847 | 0.4453 | 0.6113 | 0.8389

s2_lora | 9 | 6.3674 | 0.4735 | 0.366 | 0.418 | 0.4302 | 0.578 | 0.848

s2_lora | 10 | 6.2041 | 0.4817 | 0.3804 | 0.4378 | 0.4425 | 0.5873 | 0.8393

s2_lora | 11 | 6.0838 | 0.5103 | 0.391 | 0.4345 | 0.4593 | 0.6153 | 0.8394

s2_lora | 12 | 6.1054 | 0.4927 | 0.3869 | 0.4393 | 0.4503 | 0.6027 | 0.8395

s3_unfreeze | 1 | 6.0046 | 0.4476 | 0.3513 | 0.4431 | 0.4178 | 0.544 | 0.8389

s3_unfreeze | 2 | 5.873 | 0.4795 | 0.3702 | 0.4452 | 0.4398 | 0.6 | 0.8451

s3_unfreeze | 3 | 5.5191 | 0.5062 | 0.3751 | 0.4359 | 0.4528 | 0.5847 | 0.8555

s3_unfreeze | 4 | 5.3217 | 0.5494 | 0.3927 | 0.4157 | 0.4756 | 0.632 | 0.8632

s3_unfreeze | 5 | 5.1688 | 0.5772 | 0.3994 | 0.3945 | 0.4873 | 0.6747 | 0.8622

s3_unfreeze | 6 | 5.0934 | 0.6065 | 0.4119 | 0.3893 | 0.5047 | 0.6967 | 0.873

s3_unfreeze | 7 | 4.8659 | 0.5569 | 0.4008 | 0.4203 | 0.4827 | 0.6453 | 0.8704

s3_unfreeze | 8 | 4.8351 | 0.5841 | 0.4212 | 0.4017 | 0.4988 | 0.6633 | 0.871

s3_unfreeze | 9 | 4.7282 | 0.5883 | 0.4197 | 0.4173 | 0.5035 | 0.672 | 0.8698

s3_unfreeze | 10 | 4.6162 | 0.5293 | 0.4016 | 0.4659 | 0.4783 | 0.5853 | 0.8646

s3_unfreeze | 11 | 4.5865 | 0.5898 | 0.4272 | 0.4342 | 0.5099 | 0.67 | 0.8693

s3_unfreeze | 12 | 4.6462 | 0.599 | 0.4186 | 0.3984 | 0.5047 | 0.6827 | 0.8725

s3_unfreeze | 13 | 4.5991 | 0.5958 | 0.4321 | 0.4285 | 0.5132 | 0.6727 | 0.8713

s3_unfreeze | 14 | 4.7401 | 0.5786 | 0.4304 | 0.4417 | 0.5067 | 0.6613 | 0.8712

checkpoints saved this run:

stage | epoch

1 | 1

1 | 2

1 | 3

1 | 4

1 | 5

2 | 1

2 | 2

2 | 3

2 | 4

2 | 5

2 | 6

2 | 7

2 | 8

2 | 9

2 | 10

2 | 11

2 | 12

3 | 1

3 | 2

3 | 3

3 | 4

3 | 5

3 | 6

3 | 7

3 | 8

3 | 9

3 | 10

3 | 11

3 | 12

3 | 13

last live step:

stage: 4

epoch in stage: 1

step: 950

total steps: 3000

train loss: 4.1472

throughput img s: 2.36

monitoring time series file: data/training_monitor.csv (1160 rows so far)

learning rate per step: NOT APPLICABLE - not implemented in this codebase (not logged per-step; per-stage LRs are in section 6 config)

14 final dr grading metrics: PENDING - computed after training/evaluation completes (test-set accuracy/balanced-accuracy/macro-F1/weighted-F1/QWK/MAE, per-grade P/R/F1/support, confusion matrix, calibration NLL/ECE/Brier - all from scripts/04_evaluate.py, which runs once after training)

15 raw prediction outputs: PENDING - computed after training/evaluation completes (per-image predictions/probabilities are written by scripts/05_predict.py per image and by scripts/04_evaluate.py in aggregate - not yet run against a finished checkpoint)

16 explainability xai: PENDING - computed after training/evaluation completes (attribution AUROC/Dice/IoU vs the real 838 IDRiD/DDR masks, counterfactual sensitivity - scripts/21_xai_evaluation.py, post-training)

17 uncertainty: PENDING - computed after training/evaluation completes (MC-dropout epistemic uncertainty, decision-gate accept/reject counts on real predictions - needs a finished checkpoint; the self-test already verifies the MECHANISM works on synthetic data)

18 ablation experiments: NOT APPLICABLE - not implemented in this codebase for the full grid requested (Baseline+PI-CORAL/+DQK/+CB-Focal/+EMD/+margin/etc as SEPARATE trained arms) - this run trains ONE configuration with all components on. scripts/07_compare_models.py DOES provide baseline-vs-proposed comparisons (CNN/ Transformer baselines vs this architecture) and is part of the post-training pipeline, but it is not a component-by-component ablation grid.

19 multi seed statistics: NOT APPLICABLE - not implemented in this codebase (this run uses a single seed, 1337; no multi-seed variance/CI study is implemented or planned in this codebase)

20 cross domain evaluation: PARTIAL - scripts/04_evaluate.py scores the trained model against an APTOS-2019 external cache (data/cache_external) as a domain-shift check; a full IDRiD/DDR/EyePACS three-way grading comparison is NOT APPLICABLE - not implemented in this codebase

21 checkpoints reproducibility

checkpoint file: outputs/retfound_plus_laft_xai/checkpoint.pt (every completed epoch, overwritten in place, atomic write)

best checkpoint file: outputs/retfound_plus_laft_xai/best.pt (only when selection score improves)

contents of checkpoint pt: ['model state_dict', 'optimizer state_dict', 'scheduler state_dict', 'EMA shadow weights', 'hard-example-mining state (hardness/seen/boundary/lowconf arrays)', 'training history', 'best score/epoch/stale counter', 'torch/CUDA/numpy RNG state', 'GLA-LoRA ranks', 'original CLI args']

python rng state saved: no - only torch/CUDA/numpy RNG saved (the training loop does not use Python's own random module directly)

dataloader sampler state saved: no - WeightedRandomSampler's draw sequence is not itself seeded/checkpointed for exact replay; the underlying hardness/mining STATE that produces its weights IS checkpointed

dataset split reproducibility: deterministic given (meta.csv contents, seed 1337) via split_by_patient() - re-running produces the identical split

config json: outputs/retfound_plus_laft_xai/config.json (full config, reproduced above section-by-section)

command line args used: see resume_pipeline.ps1's 'Staged training' step for the exact invocation

18 common compute environment

cpu cores logical: 28

ram gb: 31.7

gpu model: NVIDIA GeForce RTX 5060

gpu vram gb: 8.5

gpu compute capability: 12.0 (sm_120, Blackwell)

cuda version: 12.8

pytorch version: 2.11.0+cu128

torchvision version: 0.26.0+cu128

timm version: 1.0.29

opencv version: 5.0.0

python version: 3.14.7

os: Windows-11-10.0.26200-SP0

git commit: 14e1cbcc9705725e5548ef8e7a9e37179747d72f

git commit short: 14e1cbc

power plan: High performance (switched from Balanced this session)